



Understanding Hardware Performance

Richard Rollins

github.com/rprollins/jacs-hardware

Overview

- Hardware Trends
- Roofline Performance Model
- Memory Hierarchy
- Benchmarks

Examples

The examples use the Google Benchmark library (github.com/google/benchmark).

To build the examples you should clone the repository and invoke the Makefile:

```
git clone git@github.com:rprollins/jacs-hardware.git
cd jacs-hardware
make all
```

The benchmarks were compiled with GCC 6.2.0 and run on an Intel Xeon Processor E5-1650 v4 3.60 GHz 6-core processor.

References

- Computer Architecture: A Quantitative Approach - Hennessy & Patterson
- Roofline: An Insightful Visual Performance Model for Multicore Architectures - Williams, Waterman & Patterson
- Want Fast C++? Know Your Hardware! - Timur Doumler CppCon 2016 (youtu.be/a12jYibw0vs)

Moore's Law


Fill in the blanks:

_____ doubles every _____ months?

Moore's Law and Denard Scaling


- Moore's Law: Observation that the transistor density of integrated circuits approximately doubles every 2 years.
- RC time constant is proportional to physical size, so smaller transistors can switch at higher frequencies f .
- Transistor Power Consumption: $P \approx CV^2f + I_LV$
- Denard Scaling: Transistor power density is "constant". Since 2005 this is no longer true; below 65nm non-linear leakage currents (I_L) and threshold voltage lead to a "power wall".
- New processors are not faster, but more parallel.

Hardware Trends



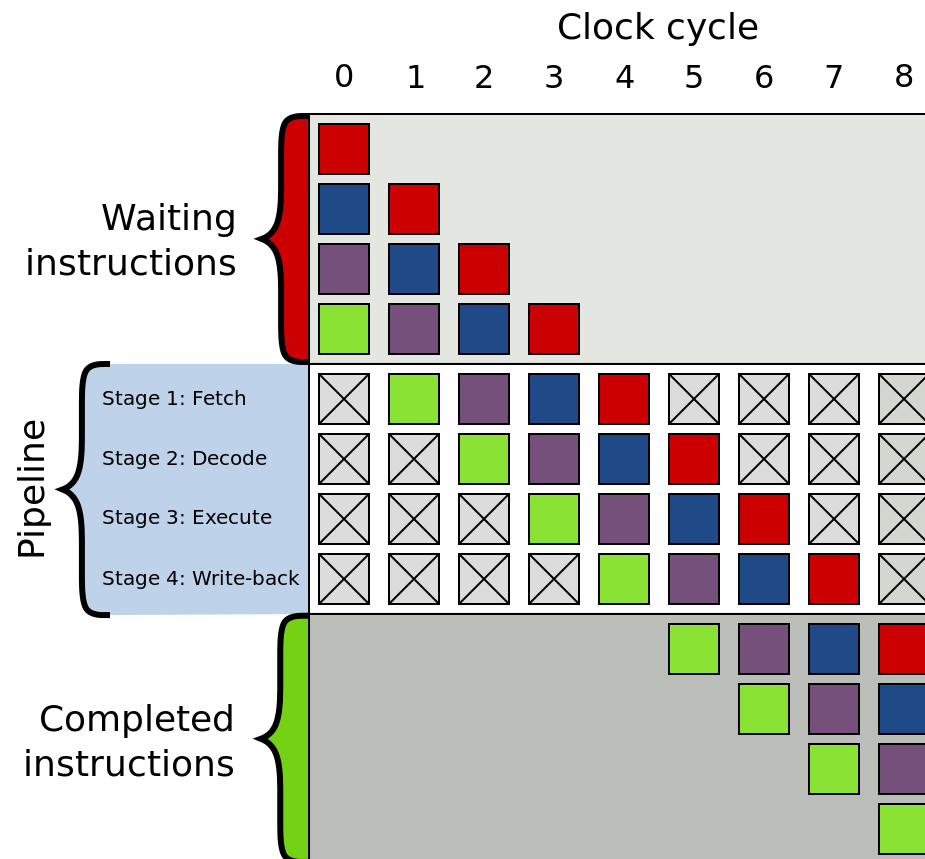
Original data collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond and C. Batten
Dotted line extrapolations by C. Moore

CPU Die




- Cores: Processor registers & local caches
- Memory Controller: Access to DRAM
- Access to peripherals (storage, network etc)

Instruction Level Parallelism - Pipelining



Instruction Level Parallelism - SIMD



- SISD: Single instruction stream, single data stream
- SIMD: Single instruction stream, multiple data streams

Example 1: Stream Benchmark

A set of benchmarks for the memory bandwidth of vector operations (www.cs.virginia.edu/stream/). Which is faster?

```
size_t N = 4096;
std::vector<double> a(N), b(N), c(N);
double q;


// BM_Add
for(size_t i=0; i<N; i++) a[i] = b[i] + c[i];
// BM_Triad
for(size_t i=0; i<N; i++) a[i] = q * b[i] + c[i];
```

Example 1: Stream Results

Benchmark	Time	Iterations
BM_Add	1441 ns	352206
BM_Triad	1404 ns	473725

How can we understand the observed performance in terms of the computer architecture?

A Model for Computation, Communication and Locality



Performance measured in FLOPS: Floating-Point Operations Per Second

Roofline Performance Model

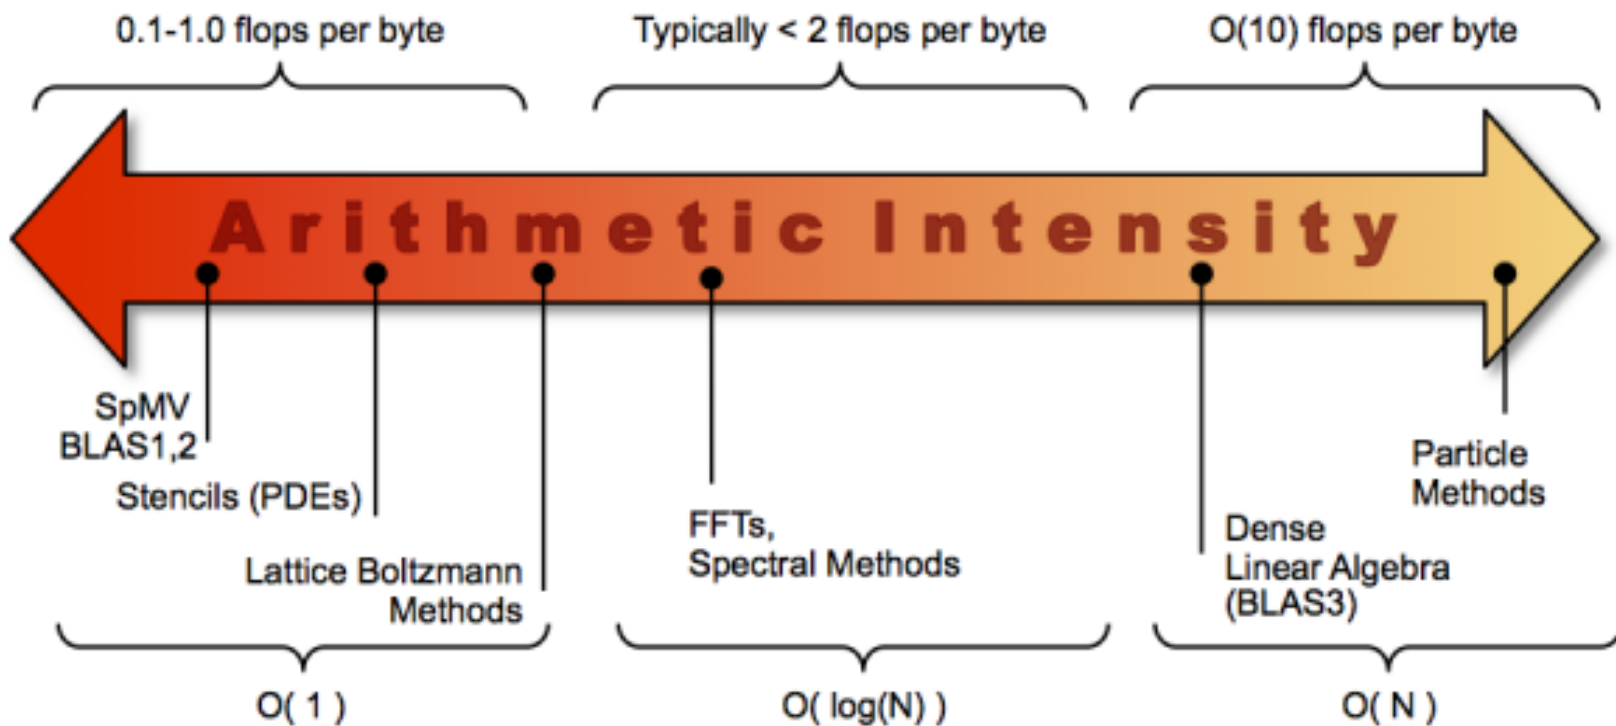
Peak performance falls into two regimes depending on the arithmetic intensity: communication-bound and compute-bound.

$$\text{Performance} = \min \begin{cases} \text{Peak Flop Rate} \\ \text{Bandwidth} \times \text{Flop:Byte Ratio} \end{cases}$$

crd.lbl.gov/departments/computer-science/PAR/research/roofline

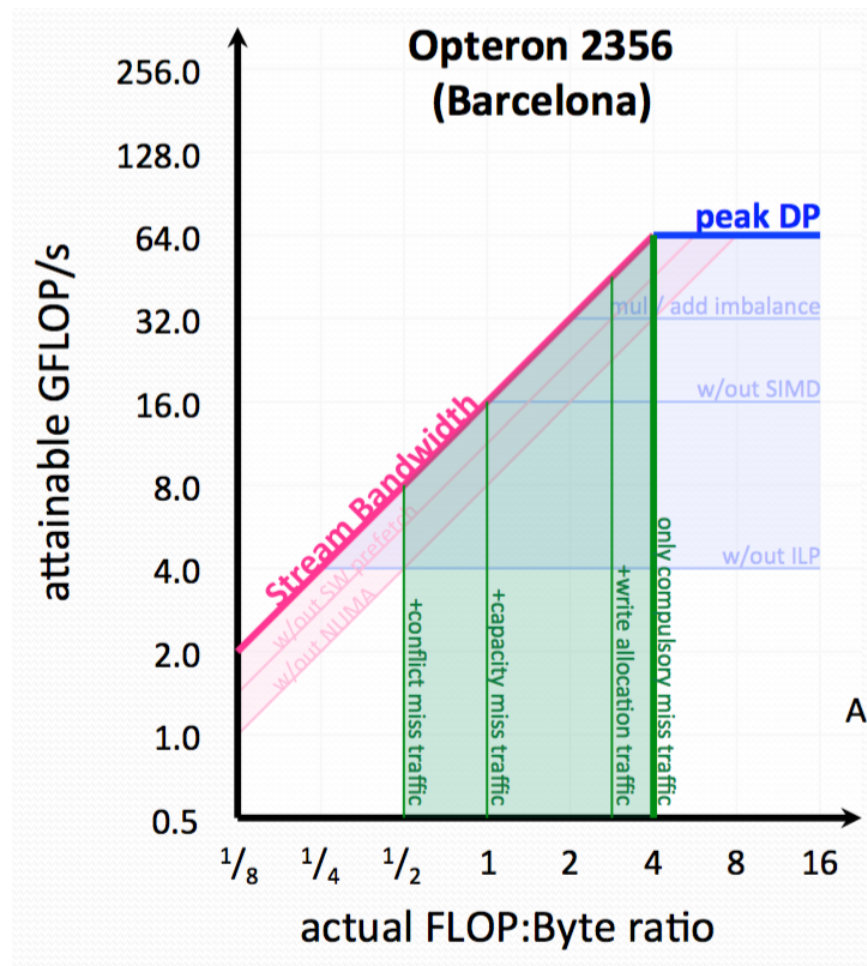
Arithmetic Intensity

A measure of data reuse for a given computer program given by the ratio of the number of floating-point operations to the volume of memory accesses.



Roofline Performance Model


Three regimes for performance optimisation: in-core, memory throughput and data reuse:



Example 1: Stream Benchmark

- How long to read data? $\text{Data Volume} / \text{Memory Bandwidth}$.
- How long to do arithmetic? $\text{Flop Count} / \text{Clock Frequency}$.
- The time to complete the arithmetic operations is negligible compared to the time to read the arrays from memory.
- Arithmetic intensity is low so the loops are memory bound.

The Processor - Memory Gap

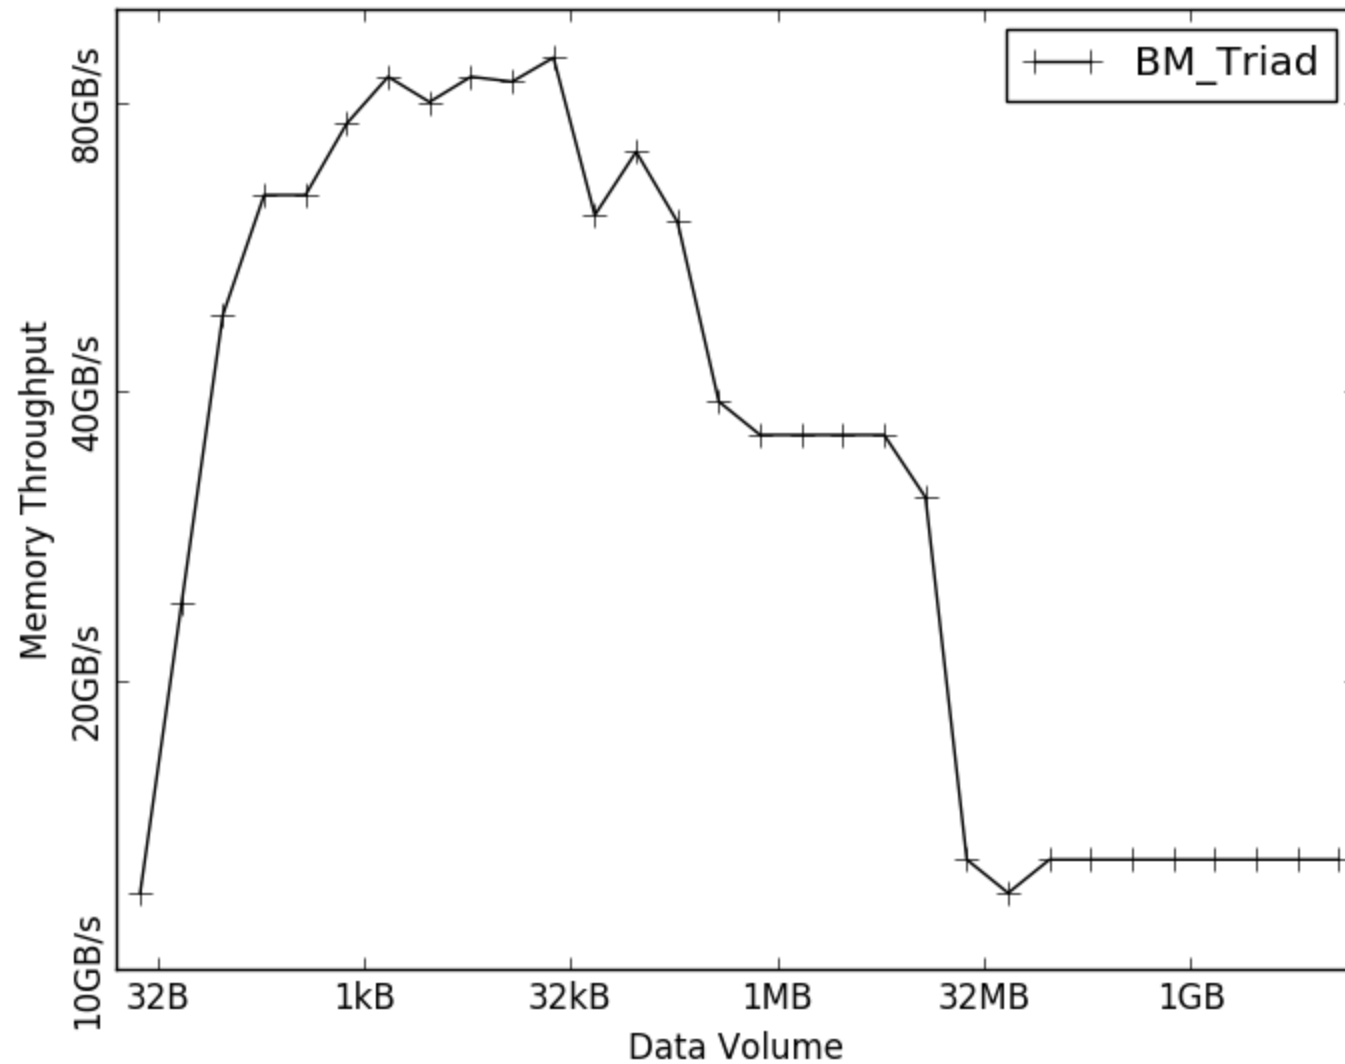


Example 2: Array Size

```
// BM_Triad  
size_t N;  
std::vector<double> a(N), b(N), c(N);  
double q;  
for(size_t i=0; i<N; i++) a[i] = q * b[i] + c[i];
```

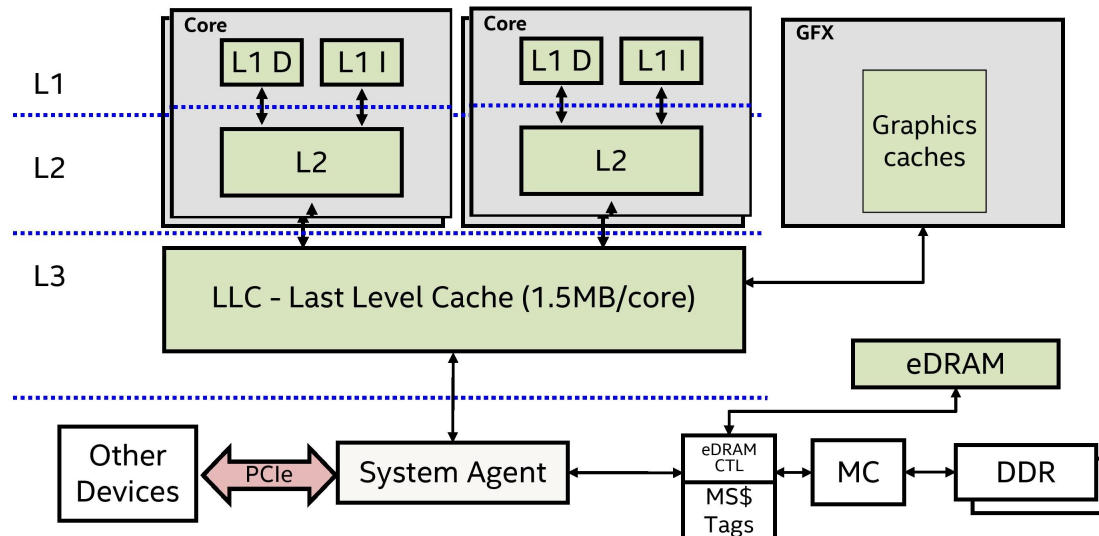
How does runtime depend on the array size N?

Example 2: Array Size




Memory Heirachy

Skylake-H (Core i7-6770HQ) Memory Hierarchy




- Memory "closer" to the processor tends to have higher bandwidth and lower latency for read/write access.
- Memory "further" from the processor tends to have greater capacity and may be non-volatile.

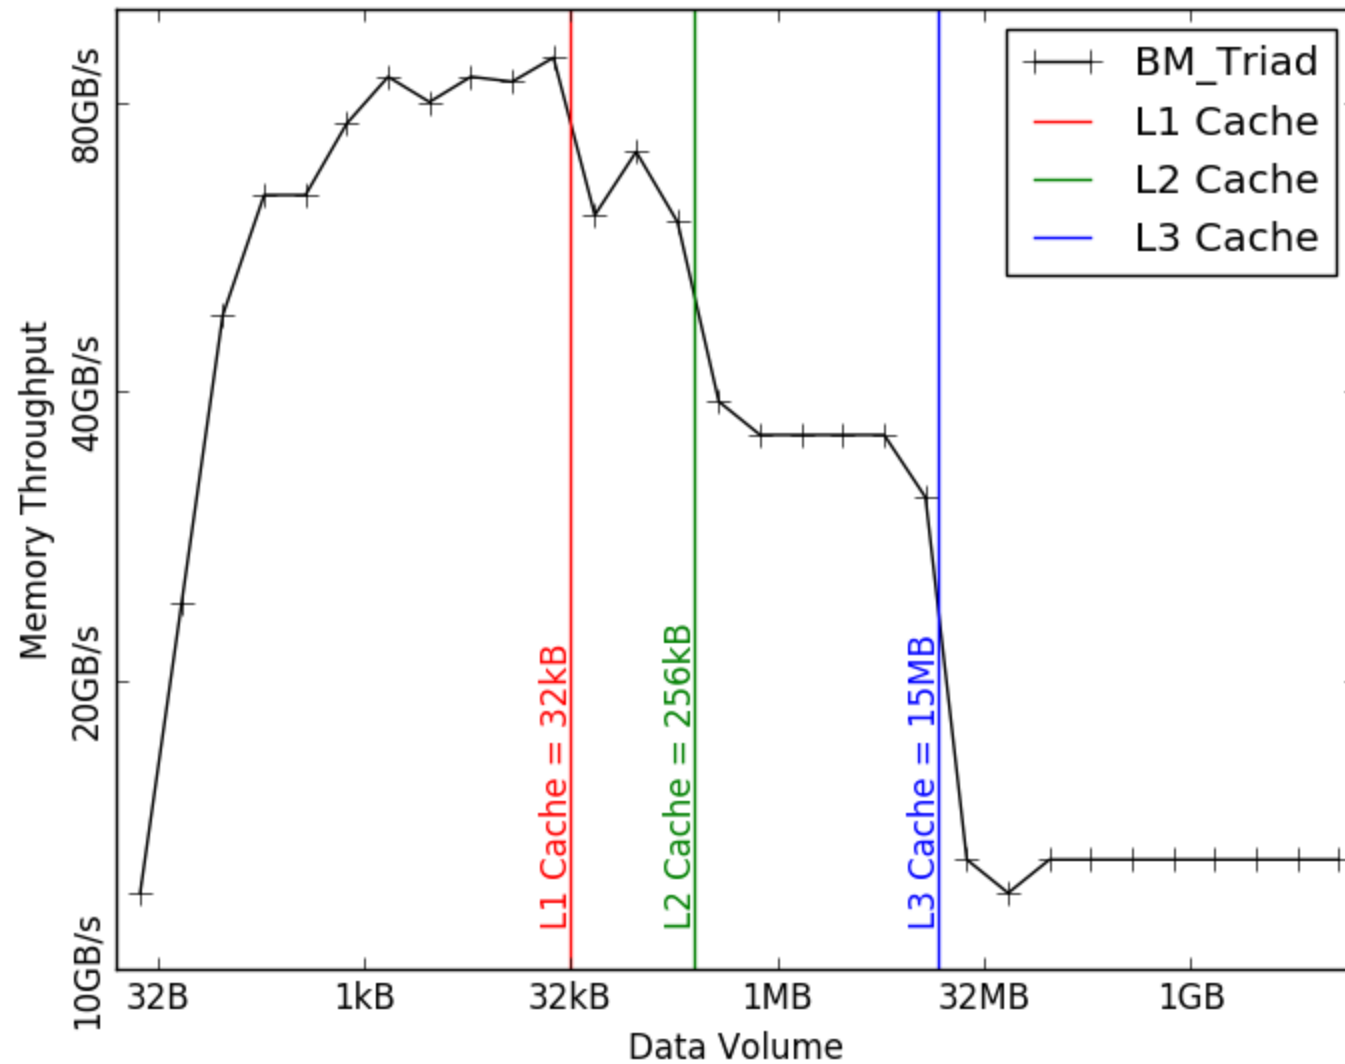
Memory Heirachy



Example 2: Array Size



Example 2: Array Size

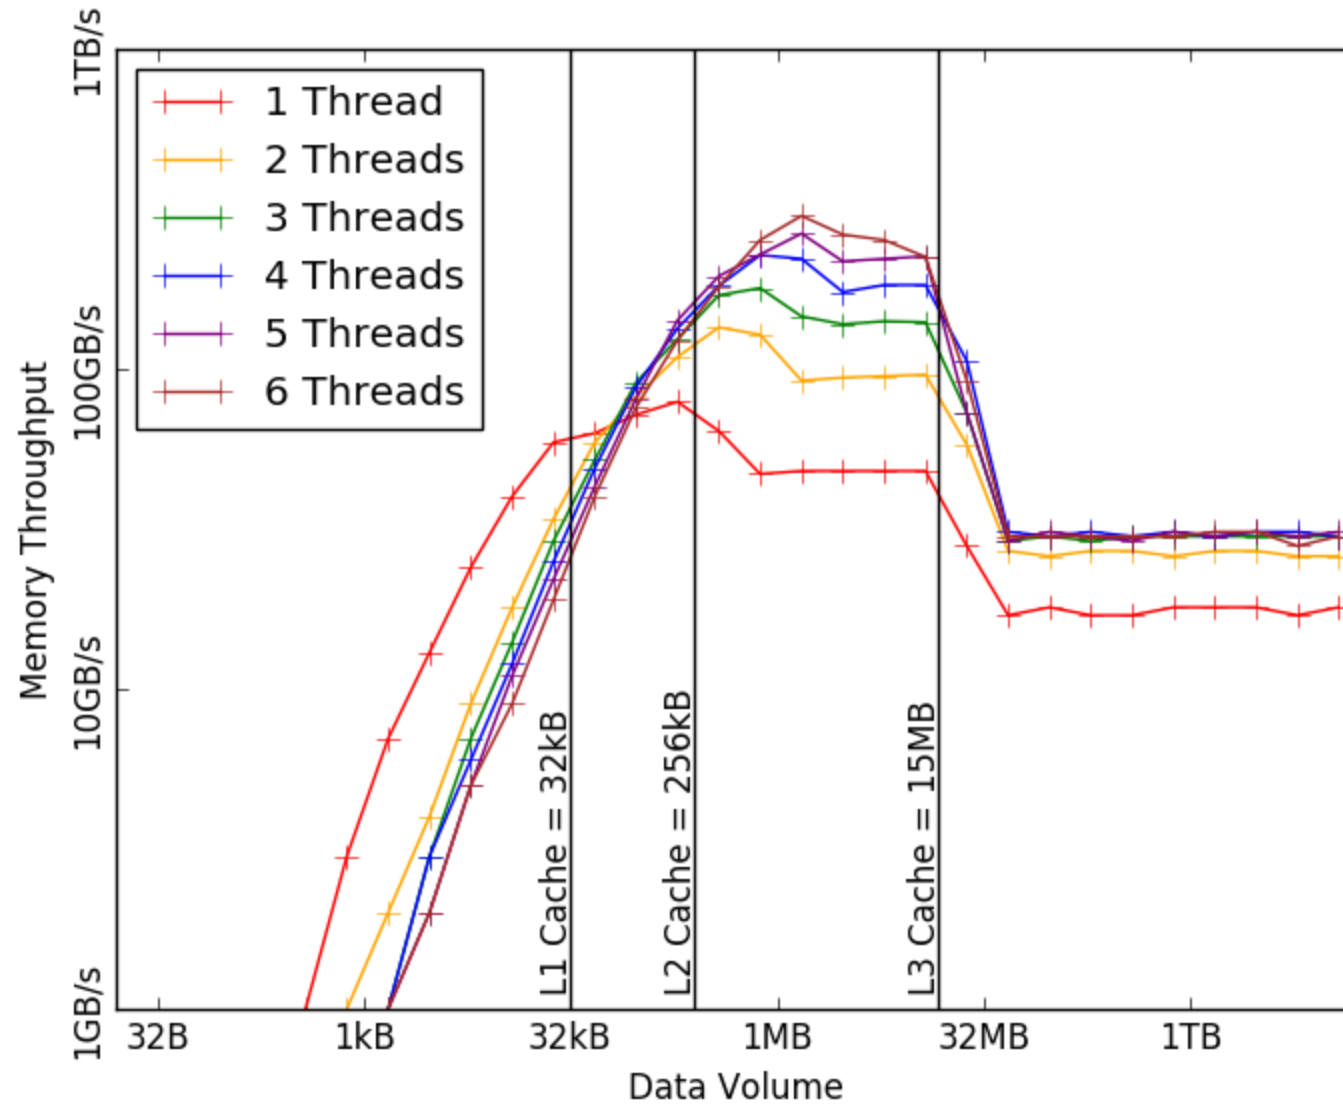


Example 3: OpenMP Stream

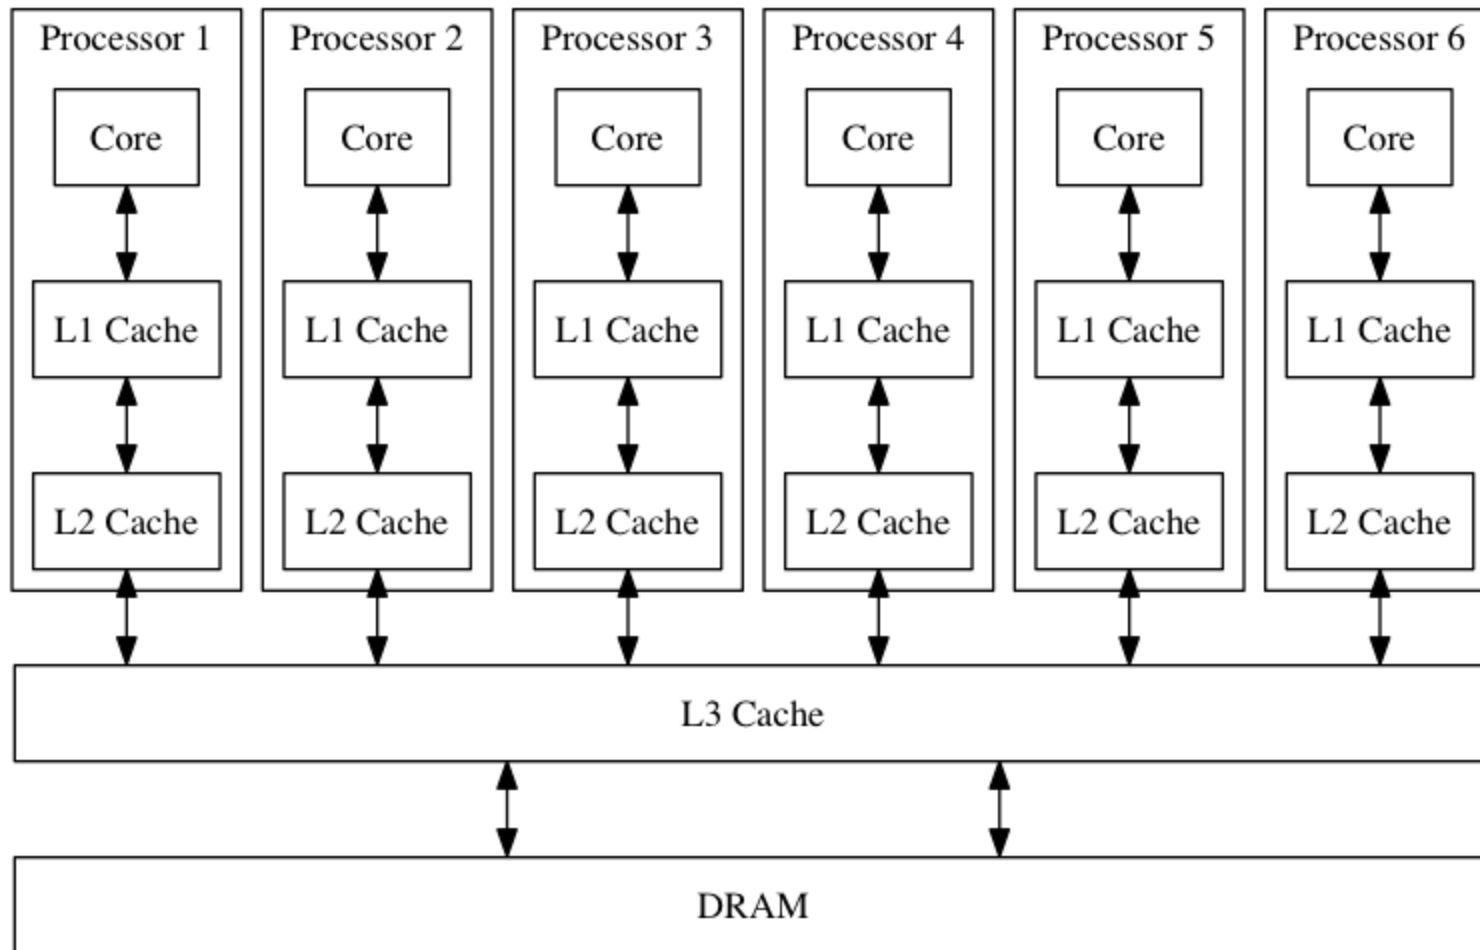
```
// BM_Triad_OpenMP  
size_t N;  
std::vector<double> a(N), b(N), c(N);  
double q;  
#pragma omp parallel for  
for(size_t i=0; i<N; i++) a[i] = q * b[i] + c[i];
```

OpenMP: shared-memory multiprocessing of loops and tasks.
How does parallelising the scale benchmark affect runtime?

Example 3: OpenMP Stream



Multicore Memory Hierarchy



Example 4: Vector vs List


```
// BM_Vector  
std::vector<double> a(N), b(N);  
std::copy(a.begin(), a.end(), b.begin());  
  
// BM_List  
std::list<double> c(N), d(N);  
std::copy(c.begin(), c.end(), d.begin());
```

Is the vector loop faster or slower than the list loop?

Example 4: Vector vs List

Benchmark	Time	Iterations
BM_Vector	1130 ns	595552
BM_List	8807 ns	77921

Example 4: Vector vs List



Each node in a linked list stores both data and the address of the next node in memory. The time taken to read successive nodes is limited by the memory latency.

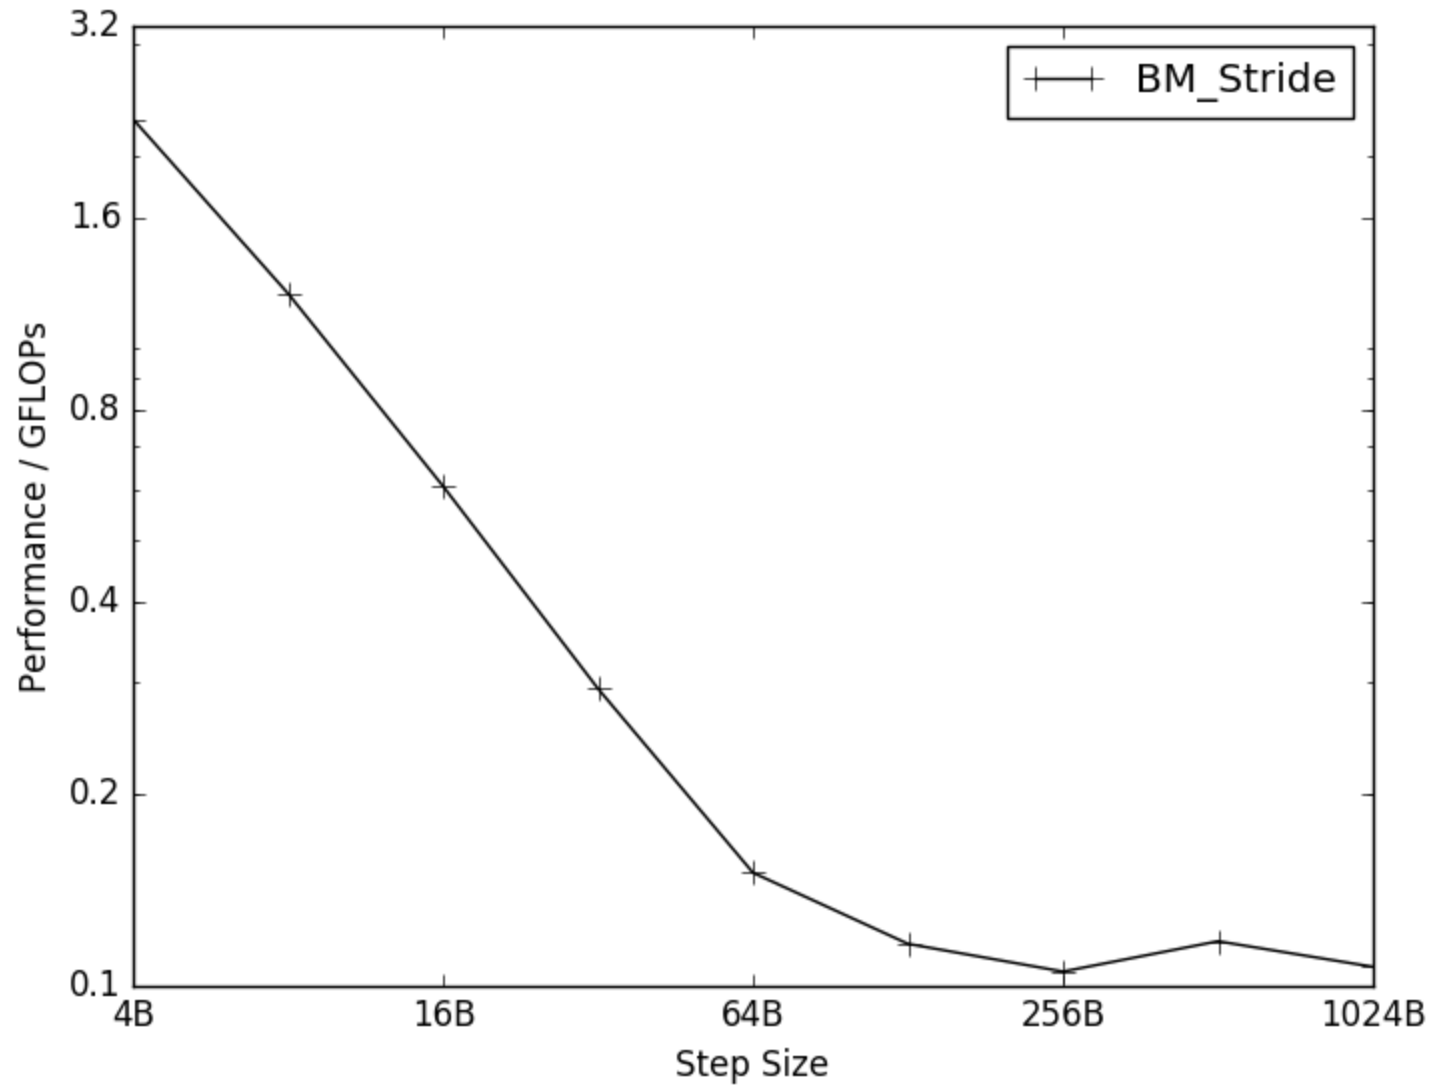
Arrays store data contiguously in memory. Modern processors *prefetch* successive elements using the full memory bandwidth.

Example 5: Stride


```
// BM_Stride  
std::vector<double> a(N); // 16MB  
for(size_t i=0; i<N; i+=stride) { a[i]++; }
```

Strided access of the data in vector a. How does the performance vary with stride size?

Example 5: Stride




Cache Lines



- Cache's cannot read in individual bits from RAM. They can only read (64B) cache lines.
- For strides less than the cache line size, all data is read.
- For strides larger than the cache line size, not all data is read.

Cache Line Invalidation and False Sharing

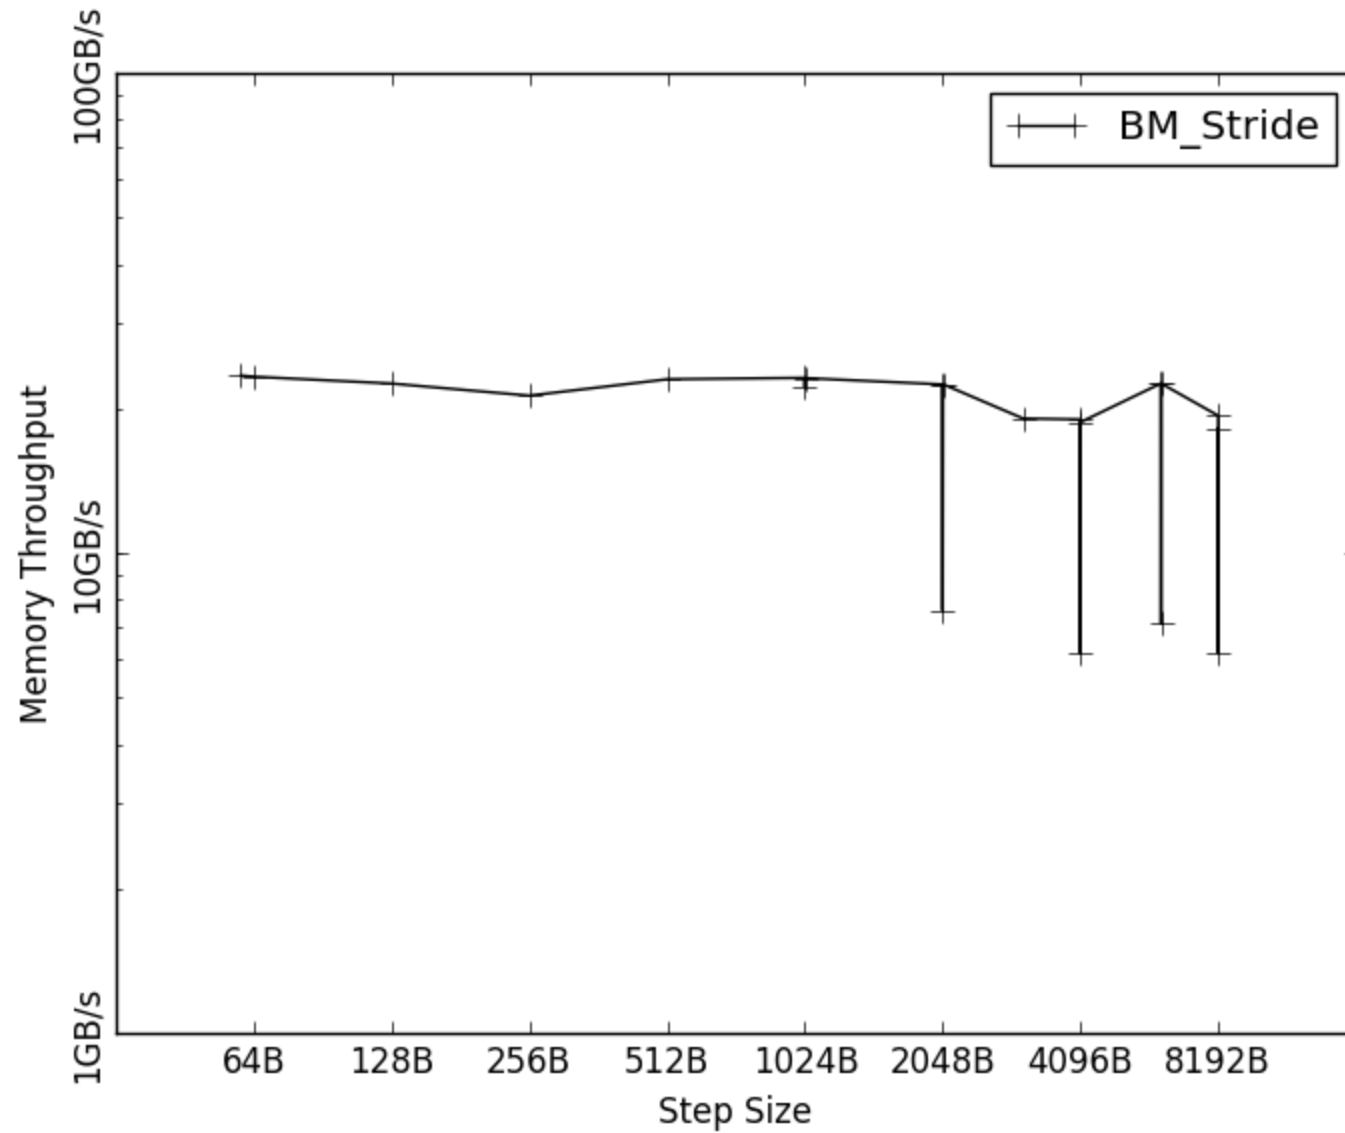


Example 6: Stride

```
// BM_Stride  
std::vector<double> a(N); // 16MB  
size_t repeat = 10000;  
size_t i = 0;  
for(size_t r=0; r<repeat; r++)  
{  
    a[i]++;  
    i = (i + stride) % N;  
}
```


Operate on a fixed number of elements in the vector *a*, but taking strides of different sizes. How does the performance vary with stride size?

Example 6: Stride




Cache Associativity

Direct Mapped
Cache Fill



Each location in main memory can be cached by just one cache location.

2-Way Associative
Cache Fill



Each location in main memory can be cached by one of two cache locations.

Memory Optimisation Strategies

- Can I shrink / divide the problem size?
- Can I use contiguous arrays?
- Can I change the data access pattern?
- Can I improve data reuse?